

E081

Azure AI Fundamentals

Armored Drill

8 concepts · 12 items · 4 transfer items

What this is. Every concept in this course's public description, re-asked with more disguise than a practice test uses. Each stem carries facts that do not decide it. Every wrong answer is a named trap.

Source. Scope from the course's public description in the WGU Institutional Catalog (September 2026); items written from the field's standard body of knowledge. No WGU assessment item, course material or proprietary content is used or reproduced. Not affiliated with or endorsed by Western Governors University.

Method © 2026 Erik Daniel Lindquist — Course Mastery Coach, armored drill. Items drafted with AI assistance under his direction.

How to run it

Paper first. Pencil. One pass, no notes, a timer if you want the pressure.

1. **Before you read the options**, fill the strip under the item: the target word (what is actually being asked), and the one fact in the stem that decides it.
2. **Principle before arithmetic**. Name the rule, then compute.
3. **Camouflage**. Most of every stem does not matter. Ask: which single fact decides this?
4. **Inverse pairs are the bait**. If two options are opposites, one of them is the trap built for you.
5. **Mark sure, unsure or guess** honestly. Sure-and-wrong is the most valuable miss on the page — it is a wrong mental model.
6. **Absolutes** (always, never, guarantee) are usually wrong. And recognizing an answer is not the same as owning it.
7. Items with a **heavy black bar** on the left are **transfer items**: the same principle, reversed or solved for a different variable.
8. Stop at the STOP page. Then grade, fill the triage grid, and tally the traps you fell for. The biggest tally is what you drill next.

The strip, filled in, looks like this:

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	_____	

Items

1.

A credit union with 40 branches in three states wants software that watches card transactions and flags ones that look unusual compared with a member's normal spending. The project is funded by a grant. Which AI workload is this?

- A. Anomaly detection
- B. Regression
- C. Conversational AI
- D. Generative AI, because banks always use it for fraud.

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

2.

A staffing agency's résumé-screening model, hosted in a modern data center, rejects qualified women at a noticeably higher rate than equally qualified men. The agency recently rebranded. Which responsible AI principle is most directly at stake?

- A. Reliability and safety
- B. Fairness
- C. Inclusiveness
- D. Privacy and security

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

3. [TRANSFER — reversed stem / different variable]

Which action most directly supports the accountability principle?

- A. Adding screen-reader support to the app.
- B. Naming an owner and a review board that must answer for the system's decisions.
- C. Letting the model make final decisions so no person is blamed.
- D. Publishing which factors drive each prediction.

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

4.

A city bike-share program, whose bikes are painted bright green, wants to predict how many bikes will be rented tomorrow using weather and day of week. Ridership doubled last year. What type of machine learning is this?

- A. Clustering
- B. Classification
- C. Classification, because every prediction picks from a fixed set.
- D. Regression

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

5. [TRANSFER — reversed stem / different variable]

A retailer has purchase histories but no customer segment labels and wants to find natural groups of similar shoppers. Which technique fits?

- A. Classification
- B. Regression
- C. Clustering
- D. Anomaly detection

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

6.

A real-estate firm, whose founder grew up in the neighborhood, trains a model to predict a home's sale price from square footage, bedrooms, and ZIP code. Its dataset has 18,000 rows. Which column is the label?

- A. Square footage
- B. Sale price
- C. The row count, 18,000
- D. All four columns together

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

7.

A parking company wants its cameras to find and draw a box around each car in a lot so it can count open spaces. License plates are visible in many images. The lot is next to a stadium. Which computer vision task fits?

- A. Image classification
- B. Face detection
- C. Object detection
- D. Optical character recognition

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

8. [TRANSFER — reversed stem / different variable]

A model returns one label per photo: “damaged” or “not damaged.” Which task is it performing?

- A. Semantic segmentation
- B. Object detection
- C. Image classification
- D. OCR

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

9.

A news-monitoring startup with an angry-sounding tagline wants to pull the names of people, companies, and cities from thousands of articles each day. Articles arrive in English. Which NLP capability fits?

- A. Named entity recognition
- B. Key phrase extraction
- C. Language detection
- D. Sentiment analysis

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

10.

An HR chatbot built on a large language model, launched during a busy open-enrollment season, confidently gives wrong answers about the company's own leave policy. The bot has a friendly avatar. What is the best fix?

- A. Raise the temperature so it explores more answers.
- B. Lengthen the prompt; long prompts guarantee correct answers.
- C. Train a brand-new foundation model from scratch.
- D. Ground the prompts with retrieved passages from the official policy documents.

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

11.

An accounts-payable team at a hospital group, which processes 5,000 invoices a month, wants to pull vendor name, date, and total from scanned invoices. The team works in a new wing. Which Azure option fits best?

- A. Azure AI Document Intelligence with its prebuilt invoice model
- B. Azure AI Speech
- C. Train a custom model from scratch in Azure Machine Learning
- D. Azure AI Translator

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

12. [TRANSFER — reversed stem / different variable]

A team must build, train, and deploy a custom model on its own sales data, using automated ML to compare algorithms. Which Azure service fits?

- A. Azure OpenAI Service
- B. Azure AI Vision
- C. Any prebuilt AI service, because prebuilt services always fit custom data.
- D. Azure Machine Learning

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

STOP

Answers and dissection follow. Finish every item and every strip first.

Key and dissection

Struck-through text is the camouflage: facts that do not decide the item.

1. Answer A Concept 1

A credit union ~~with 40 branches in three states~~ wants software that watches card transactions and flags ones that look unusual compared with a member's normal spending. ~~The project is funded by a grant.~~ Which AI workload is this?

Target word: workload type **Decides:** Flagging unusual patterns against normal behavior.

A	Correct	The goal is spotting departures from normal patterns.
B	Adjacent category	Regression predicts a number; the task here is flagging outliers.
C	Confusion set	No dialogue with users is involved.
D	Absolute language	Nothing is being generated, and no workload is universal to an industry.

2. Answer B Concept 2

A staffing agency's résumé-screening model, ~~hosted in a modern data center~~, rejects qualified women at a noticeably higher rate than equally qualified men. ~~The agency recently rebranded.~~ Which responsible AI principle is most directly at stake?

Target word: principle at stake **Decides:** Unequal outcomes for a group with equal qualifications.

A	Adjacent category	Concerns consistent, safe operation, not group bias.
B	Correct	Fairness means similar people are treated similarly across groups.
C	Confusion set	Inclusiveness is about access for people of all abilities and backgrounds, not biased outcomes.
D	Camouflage	Reacts to the data-center detail; hosting does not create the bias.

3. Answer B Concept 2 · transfer

Which action most directly supports the accountability principle?

Target word: reverse: principle to action **Decides:** Accountability = people and the organization answer for the system.

A	Adjacent category	That is inclusiveness.
B	Correct	Accountability assigns human responsibility and governance.
C	Inverted operation	Removes human responsibility, the opposite of accountability.
D	Confusion set	That is transparency.

4. Answer D Concept 3

A city bike-share program, ~~whose bikes are painted bright green~~, wants to predict how many bikes will be rented tomorrow using weather and day of week. ~~Ridership doubled last year.~~ What type of machine learning is this?

Target word: ML type **Decides:** The output is a number (count of rentals).

A	Adjacent category	Clustering groups unlabeled data; this has a known target.
B	Confusion set	Classification predicts a category, not a count.
C	Absolute language	A count is a numeric value, not a fixed label.
D	Correct	Predicting a numeric quantity is regression.

5. Answer C Concept 3 • transfer

A retailer has purchase histories but no customer segment labels and wants to find natural groups of similar shoppers. Which technique fits?

Target word: ML type without labels **Decides:** No labels exist; the goal is grouping.

A	Dropped element	Needs predefined labels, which the stem says do not exist.
B	Adjacent category	Predicts a number, not groups.
C	Correct	Unsupervised grouping of similar records.
D	Confusion set	Finds outliers, not natural groups.

6. Answer B Concept 4

A real-estate firm, ~~whose founder grew up in the neighborhood,~~ trains a model to predict a home's sale price from square footage, bedrooms, and ZIP code. ~~Its dataset has 18,000 rows.~~ Which column is the label?

Target word: identify the label **Decides:** The model predicts sale price.

A	Confusion set	An input feature.
B	Correct	The label is the value being predicted.
C	Wrong stem number	The dataset size is not a column.
D	Absolute language	Only the predicted value is the label.

7. Answer C Concept 5

A parking company wants its cameras to find and draw a box around each car in a lot so it can count open spaces. ~~License plates are visible in many images. The lot is next to a stadium.~~ Which computer vision task fits?

Target word: vision task **Decides:** Must locate and count each object.

A	Adjacent category	Gives one label for the whole image; cannot count or locate cars.
B	Confusion set	Detects faces, not vehicles.
C	Correct	Detection returns a label and a bounding box for each object.
D	Camouflage	Reacts to the license-plate detail; reading text does not count cars.

8. Answer C Concept 5 • transfer

A model returns one label per photo: "damaged" or "not damaged." Which task is it performing?

Target word: reverse: output to task **Decides:** One label for the whole image, no locations.

A	Adjacent category	Segmentation labels every pixel.
B	Inverted operation	Detection would return boxes around each item.
C	Correct	A single label per image is classification.
D	Confusion set	Reads text; no text is involved.

9. Answer A Concept 6

A news-monitoring startup ~~with an angry-sounding tagline~~ wants to pull the names of people, companies, and cities from thousands of articles each day. ~~Articles arrive in English.~~ Which NLP capability fits?

Target word: NLP capability **Decides:** Typed names of people, organizations, and places.

A	Correct	Finds and types entities such as people, organizations, and locations.
B	Adjacent category	Returns main topics, not typed named entities.
C	Confusion set	Names the language, which is already known.
D	Camouflage	Reacts to the tagline; opinion is not the goal.

10. Answer D Concept 7

An HR chatbot built on a large language model, ~~launched during a busy open-enrollment season,~~ confidently gives wrong answers about the company's own leave policy. ~~The bot has a friendly avatar.~~ What is the best fix?

Target word: reduce made-up answers **Decides:** The model lacks grounding in company-specific documents.

A	Inverted operation	Higher temperature adds randomness, making errors more likely.
B	Absolute language	Length alone does not add missing facts.
C	Adjacent category	Costly and unnecessary; grounding addresses the gap.
D	Correct	Retrieval-augmented grounding gives the model trusted facts to answer from.

11. Answer A Concept 8

An accounts-payable team at a hospital group, ~~which processes 5,000 invoices a month,~~ wants to pull vendor name, date, and total from scanned invoices. ~~The team works in a new wing.~~ Which Azure option fits best?

Target word: service selection **Decides:** Common document-field extraction has a prebuilt service.

A	Correct	Extracts standard invoice fields without custom training.
B	Dropped element	Ignores that the input is scanned documents, not audio.
C	Adjacent category	Custom training is unnecessary when a prebuilt model fits.
D	Confusion set	Translates text; does not extract fields.

12. Answer D Concept 8 • transfer

A team must build, train, and deploy a custom model on its own sales data, using automated ML to compare algorithms. Which Azure service fits?

Target word: reverse: custom need to service **Decides:** Custom training on the organization's own data.

A	Adjacent category	Provides generative models, not automated ML on tabular data.
B	Confusion set	A prebuilt image service, not a custom tabular platform.
C	Absolute language	Prebuilt services cover common tasks, not every organization-specific prediction.
D	Correct	Supports custom training, automated ML, and deployment.

Triage grid

Write each item number in the box it belongs to.

	Right	Wrong
Sure	Owned.	Concept inversion — top priority. A wrong mental model; it survives rewording.
Unsure / guess	Recognition, not mastery. Re-drill with the transfer.	Not learned yet.

Trap tally

For every miss, find the letter you chose in the key and tally its trap. The biggest number is the pattern to drill.

Trap	What it does	In this drill	Your misses
Dropped element	Leaves out a required term or step.	2	
Wrong stem number	A figure copied from the prompt that is not the answer.	1	
Adjacent category	The right idea in the wrong bucket.	11	
Inverted operation	Direction flipped: add/subtract, cause/effect, higher/lower.	3	
Absolute language	Always, never, guarantee — where judgment applies.	5	
Confusion set	Every option is real; the test is exact role-matching.	11	
Camouflage	Answers a stem fact that was never the question.	3	

Anchor rows — the concept list

One row per concept: the rule to lock in, the trap that catches people, and the tell in the question that signals it.

#	Concept	Rule	Trap	Tell
1	Matching scenarios to AI workloads	Name the workload by the output: unusual-pattern flags = anomaly detection; images = computer vision; text or speech meaning = NLP; dialogue = conversational AI; new content = generative AI.	Naming the workload by the industry or data source instead of the task.	A business scenario with one clear desired output.
2	Responsible AI principles	Microsoft's six principles: fairness, reliability and safety, privacy and security, inclusiveness, transparency, accountability — each addresses a distinct harm.	Mixing up fairness (equal treatment of groups) with inclusiveness (usable by people of all abilities) or transparency (explainability).	A scenario describing who is harmed or what is hidden.
3	Regression vs. classification vs. clustering	Regression predicts a number; classification predicts a known category; clustering groups unlabeled data.	Calling any prediction "classification," or using classification when no labels exist.	Is the output a quantity, a known label, or an unknown grouping?
4	Features and labels	The label is the value the model learns to predict; features are the inputs used to predict it.	Treating the most detailed input column as the label.	A dataset and a stated prediction goal.
5	Computer vision task types	Image classification labels the whole image; object detection labels and locates each object with a bounding box; OCR reads text.	Choosing classification when the task needs locations or counts.	"Where," "how many," or "bounding box" vs. one label per image.
6	NLP capabilities	Entity recognition finds named people, places, and organizations; key phrase extraction finds main topics; sentiment scores opinion; language detection names the language.	Using key phrases when the goal is typed entities.	The exact thing to pull out of the text.
7	Generative AI and grounding	Large language models predict likely text; grounding prompts with retrieved trusted data reduces made-up answers.	Assuming a larger prompt or higher creativity setting makes answers more factual.	A chatbot giving confident but wrong answers about company-specific facts.
8	Choosing the Azure service	Prebuilt Azure AI services fit common tasks (Vision, Language, Speech, Document Intelligence); Azure Machine Learning builds custom models on your data; Azure OpenAI provides generative models.	Training a custom model when a prebuilt service already does the task, or the reverse.	Whether the task is common and prebuilt or specific to the organization's own data.

Course scope (public description): This course provides learners with a foundational understanding of artificial intelligence concepts and how they are applied through Microsoft Azure services. The course emphasizes AI workloads and responsible AI, machine learning, computer vision, natural language processing, and generative AI, providing a very fundamental introduction to the foundations of each and how to implement them through Azure

services. This course is designed to prepare learners for Exam AI900 Azure AI Fundamentals Microsoft Certification.