

E121

Governance, Risk, and Compliance in the Age of Artificial Intelligence

Armored Drill

8 concepts · 12 items · 4 transfer items

What this is. Every concept in this course's public description, re-asked with more disguise than a practice test uses. Each stem carries facts that do not decide it. Every wrong answer is a named trap.

Source. Scope from the course's public description in the WGU Institutional Catalog (September 2026); items written from the field's standard body of knowledge. No WGU assessment item, course material or proprietary content is used or reproduced. Not affiliated with or endorsed by Western Governors University.

Method © 2026 Erik Daniel Lindquist — Course Mastery Coach, armored drill. Items drafted with AI assistance under his direction.

How to run it

Paper first. Pencil. One pass, no notes, a timer if you want the pressure.

1. **Before you read the options**, fill the strip under the item: the target word (what is actually being asked), and the one fact in the stem that decides it.
2. **Principle before arithmetic**. Name the rule, then compute.
3. **Camouflage**. Most of every stem does not matter. Ask: which single fact decides this?
4. **Inverse pairs are the bait**. If two options are opposites, one of them is the trap built for you.
5. **Mark sure, unsure or guess** honestly. Sure-and-wrong is the most valuable miss on the page — it is a wrong mental model.
6. **Absolutes** (always, never, guarantee) are usually wrong. And recognizing an answer is not the same as owning it.
7. Items with a **heavy black bar** on the left are **transfer items**: the same principle, reversed or solved for a different variable.
8. Stop at the STOP page. Then grade, fill the triage grid, and tally the traps you fell for. The biggest tally is what you drill next.

The strip, filled in, looks like this:

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	_____	

Items

1.

A regional hospital system with a new cancer center is expanding its use of AI in scheduling and billing. The data science lead, who has published widely, proposes setting the organization's AI risk appetite herself. Who should approve the AI risk appetite?

- A. Whoever builds each model, in every case
- B. The board or senior leadership governance body
- C. The CISO alone
- D. The data science lead

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

2.

A retailer values the customer database behind its AI recommendation engine at \$2,000,000. The CISO, who holds three certifications, estimates a breach would destroy 30% of that value and expects one every two years. What is the annualized loss expectancy?

- A. \$300,000
- B. \$1,000,000
- C. \$600,000
- D. \$1,200,000

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

3. [TRANSFER — reversed stem / different variable]

A risk has an SLE of \$400,000 and an ALE of \$100,000. What is its ARO?

- A. 4
- B. 100,000
- C. 0.25 (about once every four years)
- D. \$500,000

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

4.

A bank with a well-known sports sponsorship decides to shut down a planned facial-recognition feature for branch entry entirely after a legal review. The feature's pilot cost \$250,000. Which risk treatment is this?

- A. Transfer
- B. Acceptance, because the pilot money is already spent
- C. Avoidance
- D. Mitigation

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

5. [TRANSFER — reversed stem / different variable]

A risk falls within appetite, and controls would cost more than the expected loss. Leadership documents the decision and signs off. Which treatment is this?

- A. Transfer
- B. Acceptance
- C. Ignoring it, since acceptance never needs documentation
- D. Avoidance

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

6.

A logistics company with a fleet of electric trucks is adopting the NIST AI Risk Management Framework. A cross-functional team, meeting in a newly renovated room, documents an AI routing tool's intended use, users, and possible impacts on drivers and customers. Which AI RMF function is this?

- A. Manage
- B. Map
- C. Measure
- D. Identify

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

7.

A staffing firm with offices in Paris and Berlin plans to use an AI system to rank job applicants. The system was built by a vendor that also sponsors a cycling team. Under the EU AI Act's risk-based approach, how is this use most likely classified?

- A. Prohibited outright, like all AI used on people.
- B. Minimal risk, with no specific obligations.
- C. Limited risk, needing only a transparency notice.
- D. High-risk, with duties such as risk management, data governance, and human oversight.

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

8. [TRANSFER — reversed stem / different variable]

Under the EU AI Act, a customer-service chatbot must tell users they are talking with an AI but carries few other duties. Which tier fits?

- A. Minimal risk
- B. Prohibited
- C. High-risk
- D. Limited risk

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

9.

A product team at an insurer with a gleaming new headquarters wants to use a pretrained language model from a public hub that shows 2 million downloads. The model will summarize claims. What governance step comes first?

- A. Adopt it; 2 million downloads shows it is trustworthy.
- B. Scan only the application's code libraries.
- C. Verify provenance and integrity, review license and documentation, and add it to the AI inventory.
- D. Adopt it; open-source models are always safe.

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

10.

A food-delivery app with a cheerful mascot retrains its review classifier weekly on public user feedback. Analysts notice coordinated fake feedback designed to make the model misjudge certain restaurants. The app is popular with students. What threat category is this, and what defensive control fits?

- A. Data poisoning; validate and screen training data and review it before each retraining.
- B. Prompt injection; add stronger system prompts.
- C. A denial-of-service attack; add more servers.
- D. Model extraction; rate-limit the public API.

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

11.

At an engineering firm with a famous bridge design, staff are pasting confidential client drawings into free public AI chatbots. The firm has 900 employees. What is the best program response?

- A. Block the sites at the firewall with no policy or training.
- B. Let each department decide on its own.
- C. Ban all AI with no approved alternative.
- D. Issue an AI acceptable-use policy, provide an approved enterprise tool, train staff, and monitor for data loss.

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

12. [TRANSFER — reversed stem / different variable]

Which document type gives the step-by-step instructions for requesting approval of a new AI tool?

- A. Procedure
- B. Guideline, which is always mandatory
- C. Policy
- D. Standard

Target word	The one fact that decides	Answer	[] sure [] unsure [] guess
_____	_____	___	

STOP

Answers and dissection follow. Finish every item and every strip first.

Key and dissection

Struck-through text is the camouflage: facts that do not decide the item.

1. Answer B Concept 1

A regional hospital system ~~with a new cancer center~~ is expanding its use of AI in scheduling and billing. The data science lead, ~~who has published widely~~, proposes setting the organization's AI risk appetite herself. Who should approve the AI risk appetite?

Target word: governance authority **Decides:** Risk appetite is a governance decision.

A	Absolute language	Builders implement controls; they do not set appetite.
B	Correct	Leadership sets appetite and accountability; teams operate within it.
C	Confusion set	The CISO advises on security risk but does not set enterprise appetite alone.
D	Adjacent category	A management and technical role, not the governing body.

2. Answer A Concept 2

A retailer values the customer database behind its AI recommendation engine at \$2,000,000. The CISO, ~~who holds three certifications~~, estimates a breach would destroy 30% of that value and expects one every two years. What is the annualized loss expectancy?

Target word: ALE **Decides:** $SLE = \$2M \times 0.30 = \$600,000$; $ARO = 0.5$.

A	Correct	$\$600,000 \times 0.5 = \$300,000$.
B	Adjacent category	Asset value $\times$ ARO skips the exposure factor.
C	Dropped element	That is SLE; the yearly rate is left out.
D	Inverted operation	Divides SLE by ARO instead of multiplying.

3. Answer C Concept 2 · transfer

A risk has an SLE of \$400,000 and an ALE of \$100,000. What is its ARO?

Target word: solve for ARO **Decides:** $ARO = ALE \div SLE$.

A	Inverted operation	$SLE \div ALE$ flips the ratio.
B	Wrong stem number	Repeats the ALE.
C	Correct	$\$100,000 \div \$400,000 = 0.25$.
D	Adjacent category	Adds the two figures; ARO is a rate, not a dollar amount.

4. Answer C Concept 3

A bank ~~with a well-known sports sponsorship~~ decides to shut down a planned facial-recognition feature for branch entry entirely after a legal review. ~~The feature's pilot cost \$250,000.~~ Which risk treatment is this?

Target word: treatment type **Decides:** The activity is stopped altogether.

A	Confusion set	Would shift financial impact, for example through insurance.
B	Camouflage	Sunk cost does not define the treatment.
C	Correct	Stopping the activity removes the risk.
D	Adjacent category	Would reduce risk while continuing the feature.

5. Answer B Concept 3 • transfer

A risk falls within appetite, and controls would cost more than the expected loss. Leadership documents the decision and signs off. Which treatment is this?

Target word: reverse: description to treatment **Decides:** Formal, documented decision to keep the risk.

A	Confusion set	Nothing is shifted to another party.
B	Correct	A documented, approved decision to retain the risk.
C	Absolute language	Acceptance requires a documented owner and approval.
D	Inverted operation	Would stop the activity instead of keeping the risk.

6. Answer B Concept 4

A logistics company ~~with a fleet of electric trucks~~ is adopting the NIST AI Risk Management Framework. A cross-functional team, ~~meeting in a newly renovated room~~, documents an AI routing tool's intended use, users, and possible impacts on drivers and customers. Which AI RMF function is this?

Target word: framework function **Decides:** Establishing context and impacts.

A	Adjacent category	Manage prioritizes and treats risks after they are mapped and measured.
B	Correct	Map establishes context, intended use, and potential impacts.
C	Adjacent category	Measure assesses and tracks identified risks with metrics.
D	Confusion set	A NIST Cybersecurity Framework function, not an AI RMF function.

7. Answer D Concept 5

A staffing firm ~~with offices in Paris and Berlin~~ plans to use an AI system to rank job applicants. The system was built by a vendor ~~that also sponsors a cycling team~~. Under the EU AI Act's risk-based approach, how is this use most likely classified?

Target word: risk tier **Decides:** Employment decisions are a listed high-risk use.

A	Absolute language	Only specific practices are banned; hiring tools are regulated, not banned.
B	Dropped element	Ignores the employment context.
C	Adjacent category	Transparency-only duties fit uses such as chatbots, not hiring decisions.
D	Correct	Hiring affects people's livelihoods, a designated high-risk area.

8. Answer D Concept 5 • transfer

Under the EU AI Act, a customer-service chatbot must tell users they are talking with an AI but carries few other duties. Which tier fits?

Target word: reverse: obligations to tier **Decides:** Transparency duty only.

A	Confusion set	Minimal-risk uses have no specific transparency duty.
B	Inverted operation	Prohibited uses cannot be deployed at all.
C	Adjacent category	Would bring conformity and oversight duties.
D	Correct	Characterized by transparency obligations.

9. Answer C Concept 6

A product team at an insurer ~~with a gleaming new headquarters~~ wants to use a pretrained language model from a public hub ~~that shows 2 million downloads~~. The model will summarize claims. What governance step comes first?

Target word: supply-chain due diligence **Decides:** A third-party model is entering the environment.

A	Camouflage	Popularity is not assurance.
B	Dropped element	Misses the model and its data, the new AI components.
C	Correct	Core third-party AI due diligence before use.
D	Absolute language	Open source still needs review.

10. Answer A Concept 7

A food-delivery app ~~with a cheerful mascot~~ retrain its review classifier weekly on public user feedback. Analysts notice coordinated fake feedback designed to make the model misjudge certain restaurants. ~~The app is popular with students~~. What threat category is this, and what defensive control fits?

Target word: threat and control **Decides:** Untrusted input enters through training data.

A	Correct	The manipulation targets training data, so controls sit on the data pipeline.
B	Confusion set	Prompt injection targets run-time instructions, not training data.
C	Dropped element	Ignores that the goal is to skew the model, not overload it.
D	Adjacent category	Extraction copies a model through queries; this corrupts its training.

11. Answer D Concept 8

At an engineering firm ~~with a famous bridge design~~, staff are pasting confidential client drawings into free public AI chatbots. ~~The firm has 900 employees~~. What is the best program response?

Target word: AI program management **Decides:** Unapproved AI use exposes confidential data.

A	Adjacent category	A technical control alone leaves gaps and no guidance.
B	Dropped element	Leaves no consistent policy or accountability.
C	Absolute language	Blanket bans tend to push use underground.
D	Correct	Combines rules, a safe alternative, awareness, and oversight.

12. Answer A Concept 8 · transfer

Which document type gives the step-by-step instructions for requesting approval of a new AI tool?

Target word: reverse: content to document type **Decides:** Step-by-step instructions = procedure.

A	Correct	Procedures give the detailed steps.
B	Absolute language	Guidelines are recommended, not mandatory.
C	Adjacent category	States high-level mandatory intent, not steps.
D	Confusion set	Sets specific required values or technologies, not a workflow.

Triage grid

Write each item number in the box it belongs to.

	Right	Wrong
Sure	Owned.	Concept inversion — top priority. A wrong mental model; it survives rewording.
Unsure / guess	Recognition, not mastery. Re-drill with the transfer.	Not learned yet.

Trap tally

For every miss, find the letter you chose in the key and tally its trap. The biggest number is the pattern to drill.

Trap	What it does	In this drill	Your misses
Dropped element	Leaves out a required term or step.	5	
Wrong stem number	A figure copied from the prompt that is not the answer.	1	
Adjacent category	The right idea in the wrong bucket.	11	
Inverted operation	Direction flipped: add/subtract, cause/effect, higher/lower.	4	
Absolute language	Always, never, guarantee — where judgment applies.	6	
Confusion set	Every option is real; the test is exact role-matching.	7	
Camouflage	Answers a stem fact that was never the question.	2	

Anchor rows — the concept list

One row per concept: the rule to lock in, the trap that catches people, and the tell in the question that signals it.

#	Concept	Rule	Trap	Tell
1	Governance vs. management	The board and senior leadership set direction, risk appetite, and accountability; management and technical teams implement and operate controls within it.	Letting the team that builds AI systems set the organization's risk appetite.	"Who approves" or "who is accountable" for AI risk.
2	Quantitative risk: SLE, ARO, ALE	$SLE = \text{asset value} \times \text{exposure factor}$; $ALE = SLE \times ARO$.	Stopping at SLE or skipping the exposure factor.	An asset value, a percentage of loss, and a yearly frequency.
3	Risk treatment options	Mitigate = reduce likelihood or impact; transfer = shift financial impact (insurance, contracts); avoid = stop the activity; accept = formally take the risk within appetite.	Calling insurance mitigation, or treating acceptance as ignoring the risk.	What the organization actually does about the risk.
4	AI risk frameworks	NIST AI RMF has four functions: Govern (culture and accountability), Map (context and impacts), Measure (assess and track), Manage (prioritize and treat).	Mixing AI RMF functions with NIST CSF functions or with each other.	An activity to place in a framework function.
5	Risk-based AI regulation	The EU AI Act sorts uses into prohibited, high-risk (e.g., hiring, credit, critical services — heavy obligations), limited-risk (transparency duties), and minimal-risk tiers.	Classifying by technology rather than by use and impact.	An AI use case and a question about obligations or tier.
6	AI supply chain risk	Third-party models, datasets, and AI services need due diligence: provenance and integrity checks, license and documentation review, inventory, and contract terms.	Trusting a component because it is popular or open source.	A pretrained model, public dataset, or AI vendor being adopted.
7	AI-specific threat categories and defensive controls	Data poisoning corrupts training data; prompt injection manipulates model instructions at run time; model extraction copies a model through queries — each maps to different defensive controls.	Naming the threat by the interface instead of by what is being manipulated.	Where untrusted input enters: training data, prompts, or query access.
8	AI policy and program management	Policies state mandatory intent, standards set specific requirements, procedures give steps, guidelines advise; effective AI programs pair policy with approved tools, training, and monitoring.	Relying on a blanket ban or a technical block without policy and alternatives.	Employees using unapproved AI tools, or a question about which document type.

Course scope (public description): Governance, Risk, and Compliance in the Age of Artificial Intelligence prepares learners to design cybersecurity governance frameworks that address the evolving risks introduced by AI adoption. Learners evaluate regulatory compliance requirements, analyze AI-related threats and supply chain vulnerabilities, and develop policy and program management strategies that integrate responsible AI principles to lead GRC efforts in environments where AI is both a tool and a risk surface.