

Joshua Says

Armored ServiceNow: Now Assist Skill Kit

For people who already know the easy version.

Advanced practice for the people building generative-AI skills and custom apps on the platform

by Joshua

First edition. Written from the platform's public documentation as of the date below. Check every key against the source area named in *Check the work* before relying on it; corrections are dated and kept public.

16 items · edition date 2026-09-15

ServiceNow is a trademark of its owner. This book is an independent, unofficial study aid. It is not affiliated with, sponsored by, or endorsed by ServiceNow or any employer of the author. Every item is original and written from the platform's public documentation; no certification exam question is reproduced or paraphrased, and no customer instance, configuration or data is described.

How to run it

- 1 **Read the target word first.** Write it on the line before you read the options. The platform disguises the question in the noun it uses — case or incident, skill or agent, queue or division, total or per unit.
- 2 **Name the one fact that decides.** Everything else in the stem is camouflage: the deadline, the executive, the numbers, the object's name.
- 3 **Principle before configuration.** Say what the platform construct is *for*, then pick the construct.
- 4 **Confusion sets are the platform's favourite trap.** Every wrong option is a real thing. The test is exact role-matching, not recognition.
- 5 **Mark confidence honestly.** Sure-and-wrong is a wrong mental model, and it is the only thing worth finding.
- 6 **Grade in Part B.** Each distractor carries the error it encodes. Name the mold you fell for; naming it defeats it.

Check the work — sources

Every item is written from public documentation. Before trusting a key, check it against the source area named here; corrections are dated and kept public.

ServiceNow Product Documentation — Now Assist Skill Kit	docs.servicenow.com > Now Assist > Now Assist Skill Kit (skills, prompts, tools, grounding, testing, deployment)
ServiceNow Product Documentation — Generative AI Controller	docs.servicenow.com > Generative AI Controller (providers, capabilities)
ServiceNow Product Documentation — App Engine Studio and application scope	docs.servicenow.com > App Engine (scoped applications, delegated development)

The seven molds: dropped element · wrong stem number · adjacent category · inverted operation · absolute language · confusion set · camouflage (the option that answers a stem fact that was never the question).

Part A — Items

1. What a skill is

A developer says she 'wrote a prompt' and wants it in the Incident workspace. In Skill Kit terms, what is missing before it can be used?

- A. A business rule
- B. An update set
- C. A skill wraps the prompt with inputs (grounding), an optional tool, an output definition, testing, and a deployment target; the prompt alone is one part
- D. Nothing — a prompt is a skill

Target word: _____ The one fact that decides: _____
Answer: _____ ☐ sure ☐ unsure ☐ guess

2. Skill types

Three asks: (1) shorten a long case into a paragraph; (2) write a knowledge article from a resolved incident; (3) answer questions over a set of documents. Which skill patterns?

- A. All three require a custom LLM
- B. 1 summarization · 2 generation · 3 Q&A / retrieval-grounded chat
- C. All three are summarization
- D. 1 generation · 2 summarization · 3 classification

Target word: _____ The one fact that decides: _____
Answer: _____ ☐ sure ☐ unsure ☐ guess

3. Grounding

The summary skill keeps inventing details not on the case. What is the correct fix?

- A. Lower the output length
- B. Switch providers
- C. A longer prompt telling it not to invent
- D. Ground the skill on the case record's fields and related records (comments, work notes) so the model only sees real data

Target word: _____ The one fact that decides: _____
Answer: _____ ☐ sure ☐ unsure ☐ guess

4. Tools

A skill must look up the customer's entitlement in another table before answering, and the value must be exact. Which Skill Kit component performs the lookup?

- A. A tool — a script, subflow, or retriever that fetches data deterministically and passes it to the prompt
- B. A UI action
- C. An ACL
- D. The prompt, by asking the model to recall it

Target word: _____ The one fact that decides: _____
Answer: _____ ☐ sure ☐ unsure ☐ guess

5. Generative AI Controller and providers

Security asks where the prompt goes. The skill uses 'Now LLM' — what does that mean, and what is the alternative path?

- A. There is no way to know
- B. The prompt runs on the instance's own CPU
- C. The prompt is sent to ServiceNow-hosted models via the Generative AI Controller; the alternative is a configured third-party provider (e.g., through a connector), which is a separate data path and approval
- D. The prompt goes to whichever model is cheapest

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

6. Testing a skill

Before publishing, the developer wants to run the skill against ten real-looking records and compare outputs. Where?

- A. In a spreadsheet
- B. In Skill Kit's test/evaluation area — run against sample records, review outputs, iterate on prompt and grounding
- C. Testing is not possible for generative skills
- D. In production on live cases

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

7. Publishing and deployment

The skill tests well. Users still do not see it on the incident record. What are the remaining steps?

- A. Restart the instance
- B. Add the user to admin
- C. Nothing — publishing deploys everywhere
- D. Publish the skill, then deploy it to a capability/UI surface (e.g., a workspace panel or UI action) and enable it in Now Assist admin

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

8. Roles and delegated development

A business analyst wants to build skills but must not modify ACLs or global scripts. How is this arranged?

- A. Skill Kit roles for building skills plus delegated development permissions inside a scoped application; no admin
- B. Use the global scope
- C. Have IT build everything
- D. Grant admin temporarily

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

9. Application scope

A custom app built in App Engine Studio needs to read incidents and write to its own tables. A developer gets an error writing to incident. Why?

- A. The app needs the admin role
- B. The app must be moved to global
- C. Cross-scope access is governed; a scoped app can be granted read but writing to another scope's table requires explicit cross-scope privilege
- D. Incident is read-only for everyone

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

10. Output definition

A generation skill should return a title, a body, and a confidence note, and the UI should place each in its own field. What makes that possible?

- A. Asking the model to use headings
- B. A structured output definition on the skill (schema) that the deployment maps to fields
- C. A business rule that parses text
- D. Three separate skills

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

11. Sensitive data

Cases sometimes contain card numbers in comments. The skill summarizes comments. What must the design address?

- A. Use a third-party provider
- B. Nothing; the model forgets
- C. Turn off comments
- D. Data privacy controls — masking/exclusion of sensitive fields and content before grounding, per the platform's sensitive data handling

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

12. Skill vs AI agent

A stakeholder wants 'AI that fixes the incident.' The developer has built a summarization skill. What is the distinction?

- A. A skill produces content on request; an AI agent (Agent Studio) plans and takes actions through tools with human oversight — different construct, different governance
- B. They are the same feature
- C. Agents replace skills
- D. A skill can take any action if the prompt says so

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

13. Prompt design in Skill Kit

The prompt says: 'Summarize this incident.' Outputs vary wildly in length and tone. What is the most reliable improvement?

- A. Run it twice
- B. Add 'please'
- C. Specify structure, length, audience and what to exclude, and supply the grounded fields by name
- D. Raise temperature

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

14. Usage and monitoring

After launch, the owner wants to know how often the skill is used, accepted, edited or rejected. Where does that live?

- A. The system log only
- B. Now Assist usage/analytics — adoption and feedback metrics per skill
- C. It is not tracked
- D. The incident table

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

15. Retriever-grounded Q&A

A Q&A skill must answer from the approved knowledge base only, never from the model's general knowledge. What is required?

- A. A bigger model
- B. Copying the KB into the prompt
- C. A prompt that says 'only use our KB'
- D. A retriever tool over the knowledge base as grounding, with the prompt constrained to the retrieved passages

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

16. Versioning and change

A prompt change breaks a working skill in production. What should have happened?

- A. Iterate and test a new version in sub-production, capture in an update set/app version, and promote — the skill is configuration like any other
- B. Ask the vendor
- C. Edit in production; it is only text
- D. Skills cannot be versioned

Target word: _____ The one fact that decides: _____
Answer: _____ [] sure [] unsure [] guess

STOP

Answers ahead. Finish Part A and mark every confidence box, then turn.

Part B — Key and dissection

1. What a skill is — C

Rule. Skill = prompt + grounding + tool + output + deployment.

Target word. prompt vs SKILL **Decides.** A skill is prompt + data in + action out + where it appears.

Camouflage, struck. Incident workspace

A	confusion set	A business rule
B	adjacent category	An update set
C	KEY	A skill wraps the prompt with inputs (grounding), an optional tool, an output definition, testing, and a deployment target; the prompt alone is one part
D	absolute language	Nothing — a prompt is a skill

2. Skill types — B

Rule. Summarize compresses; generate creates; Q&A retrieves then answers.

Target word. shorten · write · answer **Decides.** The verb picks the pattern.

Camouflage, struck. —

A	absolute language	All three require a custom LLM
B	KEY	1 summarization · 2 generation · 3 Q&A / retrieval-grounded chat
C	dropped element	All three are summarization
D	inverted operation	1 generation · 2 summarization · 3 classification

3. Grounding — D

Rule. Grounding = the record data the skill is allowed to see. Unknown resolves to nothing.

Target word. INVENTING details **Decides.** Hallucination is a grounding problem before it is a prompt problem.

Camouflage, struck. —

A	dropped element	Lower the output length
B	adjacent category	Switch providers
C	adjacent category	A longer prompt telling it not to invent
D	KEY	Ground the skill on the case record's fields and related records (comments, work notes) so the model only sees real data

4. Tools — A

Rule. Tools fetch; the model composes. Server decides, client renders.

Target word. must be EXACT **Decides.** Deterministic facts come from tools, not from the model.

Camouflage, struck. —

A	KEY	A tool — a script, subflow, or retriever that fetches data deterministically and passes it to the prompt
B	confusion set	A UI action
C	confusion set	An ACL
D	inverted operation	The prompt, by asking the model to recall it

5. Generative AI Controller and providers — C

Rule. Know the provider before the prompt leaves.

Target word. WHERE the prompt goes **Decides.** The controller is the switchboard; provider choice is a data-path decision.

Camouflage, struck. —

A	absolute language	There is no way to know
B	inverted operation	The prompt runs on the instance's own CPU
C	KEY	The prompt is sent to ServiceNow-hosted models via the Generative AI Controller; the alternative is a configured third-party provider (e.g., through a connector), which is a separate data path and approval
D	absolute language	The prompt goes to whichever model is cheapest

6. Testing a skill — B

Rule. Test → publish → deploy, in that order.

Target word. BEFORE publishing **Decides.** Skill Kit includes testing; publishing is a separate step.

Camouflage, struck. —

A	adjacent category	In a spreadsheet
B	KEY	In Skill Kit's test/evaluation area — run against sample records, review outputs, iterate on prompt and grounding
C	absolute language	Testing is not possible for generative skills
D	inverted operation	In production on live cases

7. Publishing and deployment — D

Rule. Publish ≠ deploy ≠ enable.

Target word. tests well · not visible **Decides.** Publish makes it available; deploy places it.

Camouflage, struck. —

A	inverted operation	Restart the instance
B	absolute language	Add the user to admin
C	absolute language	Nothing — publishing deploys everywhere
D	KEY	Publish the skill, then deploy it to a capability/UI surface (e.g., a workspace panel or UI action) and enable it in Now Assist admin

8. Roles and delegated development — A

Rule. Citizen development is safe when the scope is the fence.

Target word. build skills · NOT touch ACLs **Decides.** Scoped app + delegated development bounds what a builder can change.

Camouflage, struck. —

A	KEY	Skill Kit roles for building skills plus delegated development permissions inside a scoped application; no admin
B	inverted operation	Use the global scope
C	adjacent category	Have IT build everything
D	absolute language	Grant admin temporarily

9. Application scope — C

Rule. Scoped apps own their tables; cross-scope access is declared.

Target word. its own tables vs ANOTHER scope **Decides.** Scope is a boundary with explicit doors.

Camouflage, struck. —

A	adjacent category	The app needs the admin role
B	inverted operation	The app must be moved to global
C	KEY	Cross-scope access is governed; a scoped app can be granted read but writing to another scope's table requires explicit cross-scope privilege
D	absolute language	Incident is read-only for everyone

10. Output definition — B

Rule. Define outputs; map them.

Target word. each in its OWN field **Decides.** Structured output is defined, not parsed.

Camouflage, struck. —

A	adjacent category	Asking the model to use headings
B	KEY	A structured output definition on the skill (schema) that the deployment maps to fields
C	inverted operation	A business rule that parses text
D	dropped element	Three separate skills

11. Sensitive data — D

Rule. Mask or exclude before grounding; the red line is upstream of the prompt.

Target word. card numbers in COMMENTS **Decides.** What is grounded is what is sent.

Camouflage, struck. —

A	inverted operation	Use a third-party provider
B	absolute language	Nothing; the model forgets
C	dropped element	Turn off comments
D	KEY	Data privacy controls — masking/exclusion of sensitive fields and content before grounding, per the platform's sensitive data handling

12. Skill vs AI agent — A

Rule. Skill = output. Agent = actions. Governance scales with action.

Target word. produce vs ACT **Decides.** Content vs action is the line.

Camouflage, struck. —

A	KEY	A skill produces content on request; an AI agent (Agent Studio) plans and takes actions through tools with human oversight — different construct, different governance
B	absolute language	They are the same feature
C	absolute language	Agents replace skills
D	inverted operation	A skill can take any action if the prompt says so

13. Prompt design in Skill Kit — C

Rule. Say the shape; ground the facts.

Target word. VARY wildly **Decides.** Unspecified = variable.

Camouflage, struck. —

A	adjacent category	Run it twice
B	dropped element	Add 'please'
C	KEY	Specify structure, length, audience and what to exclude, and supply the grounded fields by name
D	inverted operation	Raise temperature

14. Usage and monitoring — B

Rule. Every human decision on a skill output is a label.

Target word. used · accepted · edited **Decides.** Skills carry telemetry.

Camouflage, struck. —

A	dropped element	The system log only
B	KEY	Now Assist usage/analytics — adoption and feedback metrics per skill
C	absolute language	It is not tracked
D	adjacent category	The incident table

15. Retriever-grounded Q&A — D

Rule. Retrieve → constrain → answer with citations.

Target word. approved KB ONLY **Decides.** Retrieval bounds the answer space.

Camouflage, struck. —

A	inverted operation	A bigger model
B	dropped element	Copying the KB into the prompt
C	adjacent category	A prompt that says 'only use our KB'
D	KEY	A retriever tool over the knowledge base as grounding, with the prompt constrained to the retrieved passages

16. Versioning and change — A

Rule. Skills follow the same transport discipline as everything else.

Target word. BROKE in production **Decides.** A skill is configuration.

Camouflage, struck. —

A	KEY	Iterate and test a new version in sub-production, capture in an update set/app version, and promote — the skill is configuration like any other
B	adjacent category	Ask the vendor
C	inverted operation	Edit in production; it is only text
D	absolute language	Skills cannot be versioned

Part C — Triage

Sort every item by the two marks you made. Write item numbers in the cells.

	Right	Wrong
Sure	Owned. Bank it.	Wrong model. Top priority. Say the rule aloud; find the construct in the documentation.
Unsure	Recognition, not mastery. Redo from the rule.	Not learned. Read the rule, then the source area.
Guess	Luck. Treat as not learned.	Not learned.

Mold tally. dropped element ____ · wrong stem number ____ · adjacent category ____ · inverted operation ____ · absolute language ____ · confusion set ____ · camouflage ____ · The largest number is the pattern to drill.

ServiceNow is a trademark of its owner. This book is an independent, unofficial study aid. It is not affiliated with, sponsored by, or endorsed by ServiceNow or any employer of the author. Every item is original and written from the platform's public documentation; no certification exam question is reproduced or paraphrased, and no customer instance, configuration or data is described.

© 2026 Joshua. From Joshua's mind — that is the signature.