

ERIK LINDQUIST

Principal Software Product Manager — ServiceNow

eriklindquist2@gmail.com · 801-580-7651 · Sandy, UT · linkedin.com/in/erik-lindquist-16a90a296

SUMMARY

Owns products that span enterprise platforms rather than living inside one. **ServiceNow** is the execution layer of a loop owned end to end — where a decision becomes work that actually happens — wired to **Genesys**, a Databricks lake, ThoughtSpot, and Banner on on-premise Oracle.

Twenty years applying rules written by someone else — statute, case law, department policy — accountable for every application of them, on a clock, with incomplete information. Then a rebuild: degree and state teaching licence (2018), software engineer (2019), product manager (2021). Now writing the frameworks other people work inside.

Delivers through influence rather than authority. The platform contract, the engineering standards, and the governance constraints currently governing this work were written, negotiated and adopted without a reporting line.

WHERE SERVICENOW SITS

Genesys transcripts and sentiment → joined in the Databricks lake → scored by three in-house AI tools → surfaced for review and completed as an observation in ServiceNow → back to the lake

One loop, closed. Unstructured voice becomes structured evidence in the lake; the transform scores it; **ThoughtSpot** puts that score in front of a human; **ServiceNow** executes what the human decides; the outcome returns to the lake as a label. Every human decision improves the next machine decision as a byproduct of work already required.

Banner I&O runs alongside it — an on-premise **Oracle** student information system supporting performance assessment, linked to the loop and deliberately kept separate. Improvement work there is done in partnership with the platform team lead.

EXPERIENCE

Senior Software Product Manager Western Governors University Oct 2024 – Present

- **ServiceNow as the execution layer.** Thirteen user groups run observations in the **UQA** app; twelve had the tool and nothing to put through it. **Genesys** sentiment (enrollment) and Genesys transcripts (QA) existed separately, neither seeing the other. Joined in the **Databricks** lake, scored by three in-house AI tools against QA's rubrics, completed as an observation in **ServiceNow**. An executed workflow now produces telemetry, not a closed ticket. Mentoring is the second of the thirteen.
- **Banner I&O.** Integration and operations on the on-premise **Oracle** student information system supporting performance assessment, in partnership with the platform team lead — linked to the loop and kept separate from it.
- **Eliminated sampling.** Enrollment call review from ~5% to 100% at flat headcount. The assignment was 10%; the answer was to stop sampling. **Genesys** transcripts land in the **Databricks** lake as raw input — unstructured voice becomes a queryable asset. The Operational AI Framework — ingest, AI transform, human decision, feedback loop — is in its second production instance, instrumenting the lifecycle from enrollment through graduation.
- **Wrote the trust contract first.** Scores sequence, never evaluate — never admissible in a performance review, only in deciding which calls a human hears. Set at launch, which is why nobody fought it.
- **Architecture as governance.** Three institutional AI capabilities already in service and already cleared — a business rules engine, machine learning models, natural language processing — with QA's existing rubrics as the scoring standard and the training signal. No new ingress, egress, processor or vendor review. Incremental cost: compute time.
- **Closed the loop with the delta.** Machine-to-human agreement 60% → 80%+ by making the gap the training signal. Every coaching evaluation doubles as a label; the queue improves with no added work.
- **Business case, and the manual half.** ~\$660K in avoided hiring over five years. Setting up one observation was manual — form by hand, **Genesys** download, paste into **ServiceNow**, then **Salesforce** cases and audit logs to verify. Automating it cut ~45% from time-to-complete and met the 10% target before any AI applied. The Salesforce verification is still automatable and still unclaimed.
- **The next increment, designed and held.** **ThoughtSpot** made the **Databricks** link clickable each morning; that click was the manual step inside an automated pipeline. The **Now Assist Skill Kit** could pull pre-queried, pre-scored records with automations drafting responses for human approval — 80% precision could have been 95% before first publish, three senior engineers out of hands-on work, eight working months to eight days. Held: the Skill Kit reads as citizen development, and the platform owner's misgivings were well-founded. In **ServiceNow** developers' hands it is the efficiency. Design intact.

Product Manager Acima / Upbound 2021 – 2024

- **Generative-AI underwriting.** Directed a generative-AI assessment and categorization tool that underwrote individual transactions granularly, expanding the range of financial products the company could offer — a system that scored so a human could decide.
- **16-person delivery.** Led a 16-person development team through an overhaul of the servicing platform across UI/UX and backend servicing processes.
- **SOX at volume.** Rebuilt the accounting methodology to SOX audit compliance across 100,000+ daily financial transactions.

Software Engineer Acima / Upbound 2019 – 2021

- Built and maintained a B2B contract origination platform serving 10,000+ daily users, accountable for performance and scalability.
- Developed RESTful APIs and third-party integrations across B2B and B2C origination; defined requirements, timelines, and delivery with cross-functional teams.

Middle School Teacher — Biology & Earth Science Alpine School District 2018

- Utah-licensed in biology and earth science. Designed assessments that measured what a student actually knew — including a rocket lab where students estimated an altitude they could never reach and reported honest error bars with it.

Police Commander Murray City Corporation 1997 – 2017

- **Trained the decision, not only made it.** Field Training Officer, then Rapid-response Training Officer — twenty years transferring judgment to other officers under real consequence. The same practice as mentoring product managers, at higher stakes.
- **Command.** Commanded a 16-person tactical team and an 11-person investigative division.

EDUCATION

MBA, IT Management — Western Governors University, expected Aug 2027

BS, Biology Education — Utah Valley University, 2018

SKILLS

Enterprise platforms ServiceNow workflow and automation · Genesys contact centre · Banner I&O on on-premise Oracle · Databricks lake ingest and modelling · ThoughtSpot visualization · Now Assist Skill Kit · Salesforce · AWS · SQL

Product strategy and roadmapping across domains · steering committees and stakeholder alignment · discovery in ambiguity · business-case and ROI modelling · evaluation design · executive narrative · agile delivery · mentoring product managers


Decision intelligence feedback-loop architecture · human-in-the-loop design · model-to-human calibration · next-best-action modelling · synthetic data generation · agent orchestration

Regulated FERPA-governed data · data classification · in-house-model architecture · security-review sequencing · provenance, licensing and disclosure discipline

Technical Python · JavaScript · Ruby on Rails · Elixir / Phoenix LiveView · D3 / SVG · PostgreSQL · SQL · Git · CI/CD · Vitest · RSpec

ONE AUTHORED FRAMEWORK — THE WGU STUDENT LIFECYCLE AS A FLYWHEEL

ONE AUTHORED FRAMEWORK — THE WGU STUDENT LIFECYCLE AS A FLYWHEEL



ANSWERS, PRE-SUPPLIED — THE QUESTIONS YOU WERE GOING TO ASK

The hard questions, answered before they're asked. My recall slips when I need it most, so I wrote it down. It's what I do with every framework: stop the outcome from depending on the moment.

Two truths and a lie.

1. My career path includes police commander, teacher, software engineer, and product manager.
2. I rebuilt an accounting methodology carrying more than 100,000 transactions a day.
3. I think Boston Dynamics robots and GPT-6 Astra are going to take my job.

The lie: #3.

If a robot and a frontier model can replace the work, then it probably wasn't product management in the first place. Product management is judgment under ambiguity: you frame the problem, write the contract, align the people, allocate the risk, and decide what happens when goals collide. The tools will get better. The judgment gets more valuable.

Five years in the title. This role asks for ten-plus, at principal level. Make the case.

Five years — I'll grant that. The posting's own "experience in lieu" clause is there for exactly this. But the title undercounts the work. Twenty years of high-stakes calls under strict rules came first, and then I rebuilt on purpose: degree, licence, engineer, product manager. Writing a governed AI framework — not just running one — is principal-level work, and it's already on its second production build. The bar counts years. The work shows judgment.

You say you "authored" the Operational AI Framework. Wasn't it handed to you?

I was handed the assignment: double call-review coverage, 5% to 10%. I wrote the framework: don't sample at all. Ingest, AI transform, human decision, feedback loop. Here's the proof. Enrollment coverage went from 5% to 100% at flat headcount. My business case put the hiring we avoided at about \$660K over five years. The customer's own estimate of the upside was about \$37M over five years.

Strip the story — isn't this just a QA efficiency tool?

No. It's one lever aimed down the whole funnel, end to end. Better feedback makes better front-line staff, and that moves what each stage owns: conversion, retention, and employment. One framework, three onboardings. Enrollment is the proof. Mentors is live. The rubric can be changed without retraining, and it has an exec champion and several business sponsors. Career Accelerators is in sea trial, on a Microsoft 365 engagement. Calibration went from 60% to 80%+, because every coaching review doubles as a training label.

AI scoring people's calls — why wasn't this resisted?

Wrote the trust contract first: scores sequence, never evaluate. Never admissible in a performance review, only in deciding which calls a human hears. Set at launch, which is why nobody fought it. The models are in-house and inside existing governance, so there were no new data crossings and no procurement runway between a decision and a running system.

Name a product decision you owned that went wrong — and what it cost.

The Skill Kit path. ThoughtSpot made the Databricks link clickable each morning; that click was the manual step inside an automated pipeline. The Now Assist Skill Kit could have pulled pre-queried, pre-scored records with automations drafting responses for human approval — 80% precision could have been 95% before first publish, three senior engineers out of hands-on work, eight working months to eight days. I designed it, costed it, and held it, because the Skill Kit reads as citizen development and the platform owner had well-founded misgivings — QA wanted to be citizen developers, and she was right that a platform where the business routes around engineering stops being a platform. (I'm one. She knows.) But ServiceNow developers aren't citizen developers; in their hands the kit is the efficiency. What it cost: eight months of three senior people's hands-on time, and a capability that shipped at 80% instead of 95%. What it needs before it ships is the guardrails written first — who may build what, reviewed by whom — the same move as the trust contract. That's the framework I'd bring on day one.

Engineer to product manager in two years — how did that really happen?

The acquisition did it. AcceptanceNow became Acima, then Upbound. The finance work I'd built as an engineer — the automatic adjustments, the accounting method under Dodd-Frank inside and SOX outside, the Canadian tax logic — became the thing Accounting, AR/AP and both audit teams needed one person to answer for. I was already answering for it. The title caught up with the ownership. I didn't become a PM by asking. I became one by being the person the system already leaned on.

AND TWO QUESTIONS I WILL BE ASKING YOU

1. The first framework you'd want me to write here — not extend, write. What's the problem no one has taken up yet, and what has made it stick?
2. The framework got adopted because the trust contract — scores sequence, never evaluate — was written at launch. As agent-based systems scale across the lifecycle and into the workflows **ServiceNow** executes, where hasn't that contract been written yet? As a WGU student I keep landing on performance assessment rubrics, and the President and the Provost are asking for this, done right. Is that a direction you'd want brought inside?