

ERIK LINDQUIST

Principal AI Product Manager — Decision Intelligence

eriklindquist2@gmail.com · 801-580-7651 · Sandy, UT · linkedin.com/in/erik-lindquist-16a90a296

SUMMARY

Builds decision intelligence as a closed loop, not as a model. Ships AI inside regulated institutions with no procurement cycle — in-house models, existing governance, no new data crossings. Defines what the system must satisfy before it exists, then proves the definition with a running instance.

Twenty years applying rules written by someone else — statute, case law, department policy — accountable for every application of them, on a clock, with incomplete information. Then a rebuild: degree and state teaching licence (2018), software engineer (2019), product manager (2021). Now writing the frameworks other people work inside.

The Operational AI Framework — ingest, AI transform, human decision, feedback loop — is in its second production instance at WGU, instrumenting the student lifecycle from first contact through graduation.

THE PIPELINE, END TO END

Genesys transcripts and sentiment → joined in the data lake → scored by three in-house AI tools → surfaced for review and completed as an observation in ServiceNow → back to the lake

One loop, closed. Unstructured voice becomes structured evidence in the lake; the transform scores it; the visualization layer puts that score in front of a human; ServiceNow executes what the human decides; the outcome returns to the lake as a label. Every human decision improves the next machine decision as a byproduct of work already required.

Banner I&O runs alongside it — an on-premise Oracle student information system supporting performance assessment, linked to the loop and deliberately kept separate. Improvement work there is done in partnership with the platform team lead.

EXPERIENCE

Senior Software Product Manager Western Governors University Oct 2024 – Present

- **Eliminated sampling.** Took enrollment call-review coverage from ~5% to 100% at flat headcount — Genesys transcripts with sentiment analysis, and QA observations performed in ServiceNow. The assignment was 10%; the answer was to stop sampling altogether. Established the Operational AI Framework — ingest, AI transform, human decision, feedback loop, machine-to-machine and unattended — which is what made it possible.
- **Ingest at the source.** Genesys call transcripts land in the data lake as the pipeline's raw input, making unstructured voice a queryable asset for the first time and removing sampling as a design constraint.
- **Decision surfaces over raw output.** Built the visualization layer that turns lake data into something a reviewer can act on, so the machine's contribution is sequencing and evidence rather than a verdict.
- **Closed the loop with the delta.** Machine-to-human scoring agreement moved from 60% to 80%+ by making the gap between machine and human scores the training signal — every coaching evaluation doubles as a label, and the queue improves without anyone doing additional work.
- **Wrote the human-machine contract first.** Machine scores are non-determinative by definition: never admissible in performance evaluation, only in sequencing which calls a human reviews. Set at launch rather than after objection, and the reason the platform was adopted rather than resisted.
- **ServiceNow as the execution layer.** Observations for thirteen user groups run in the UQA app in ServiceNow. Twelve of those groups had the tool and nothing to put through it — no transcripts, no evaluation process, no way to complete a review. Built the missing half for enrollment. Genesys sentiment and Genesys transcripts already existed and were consumed separately — sentiment by the enrollment team, transcripts by the QA coaches, neither seeing the other. Joining them in the data lake, under its existing security guarantees, produced something scorable for the first time; the three in-house AI tools scored it against the QA rubrics and the observation completes in ServiceNow. An executed workflow now produces telemetry instead of a closed ticket. The mentoring organisation is the second of the thirteen.
- **Banner I&O.** Integration and operations on the on-premise Oracle student information system supporting performance assessment, in partnership with the platform team lead — linked to the loop and kept separate from it.
- **Architecture as governance.** Composed the platform from three institutional AI capabilities already in service and already cleared by security — a business rules engine, machine learning models, and natural language processing — applying the QA team's existing rubrics as the scoring standard and the training signal. No new ingress or egress, no new processor, no vendor review, and the only incremental cost was compute time. The second instance now extends to the mentoring organization on the same architecture at no added risk surface.

- **Defined requirements instead of receiving them.** Produced the specification artifact for an AI platform vendor engagement: roughly 4,000 of the institution's own prospective-student decision nodes mapped to next-best-action on 100% synthetic data — zero disclosure exposure, and the institution setting the terms.
- **Business case and ROI.** Authored the case valuing avoided hiring at approximately \$660K over a rolling five years. Separately eliminated a 5.5-minute manual step from an ~11-minute per-call process, cutting about 45% from the time to complete an observation and meeting the original 10% coverage target on automation alone, before any AI screening applied.
- **Standards that outlive the project.** Wrote the engineering standards the team now builds under: validation is enforced on the server, untrusted integrations develop against local stand-ins until security review clears them, and point-in-time data always renders with its date. Built a capacity-recovery diagnostic that surfaces unowned, undocumented and stale work before it eats the team.

Product Manager Acima / Upbound 2021 – 2024

- **Generative-AI underwriting.** Directed a generative-AI assessment and categorization tool that underwrote individual transactions granularly, expanding the range of financial products the company could offer — a system that scored so a human could decide.
- **16-person delivery.** Led a 16-person development team through an overhaul of the servicing platform across UI/UX and backend servicing processes.
- **SOX at volume.** Rebuilt the accounting methodology to SOX audit compliance across 100,000+ daily financial transactions.

Software Engineer Acima / Upbound 2019 – 2021

- Built and maintained a B2B contract origination platform serving 10,000+ daily users, accountable for performance and scalability.
- Developed RESTful APIs and third-party integrations across B2B and B2C origination; defined requirements, timelines, and delivery with cross-functional teams.

Middle School Teacher — Biology & Earth Science Alpine School District 2018

- Utah-licensed in biology and earth science. Designed assessments that measured what a student actually knew — including a rocket lab where students estimated an altitude they could never reach and reported honest error bars with it.

Police Commander Murray City Corporation 1997 – 2017

- **Trained the decision, not only made it.** Field Training Officer, then Rapid-response Training Officer — twenty years transferring judgment to other officers under real consequence. The same practice as mentoring product managers, at higher stakes.
- **Command.** Commanded a 16-person tactical team and an 11-person investigative division.

EDUCATION

MBA, IT Management — Western Governors University, expected Aug 2027

BS, Biology Education — Utah Valley University, 2018

SKILLS

Decision intelligence feedback-loop architecture · human-in-the-loop design · model-to-human calibration · next-best-action modelling · synthetic data generation · agent orchestration

Enterprise platforms ServiceNow workflow and automation · Genesys contact centre · Banner I&O on on-premise Oracle · data lake ingest and modelling · visualization and reporting layers · AWS · SQL


Product strategy and roadmapping across domains · steering committees and stakeholder alignment · discovery in ambiguity · business-case and ROI modelling · evaluation design · executive narrative · agile delivery · mentoring product managers

Regulated FERPA-governed data · data classification · in-house-model architecture · security-review sequencing · provenance, licensing and disclosure discipline

Technical Python · JavaScript · Ruby on Rails · Elixir / Phoenix LiveView · D3 / SVG · PostgreSQL · SQL · Git · CI/CD · Vitest · RSpec

ONE AUTHORED FRAMEWORK — THE WGU STUDENT LIFECYCLE AS A FLYWHEEL

ONE AUTHORED FRAMEWORK — THE WGU STUDENT LIFECYCLE AS A FLYWHEEL



ANSWERS, PRE-SUPPLIED — THE QUESTIONS YOU WERE GOING TO ASK

The hard questions, answered before they're asked. My recall slips when I need it most, so I wrote it down. It's what I do with every framework: stop the outcome from depending on the moment.

Two truths and a lie.

1. I've used synthetic data to define thousands of decision paths before we engaged an AI vendor.
2. I think the biggest AI risk is that academic integrity gets breached and FERPA gets violated.
3. I built an AI framework now operating across several business functions.

The lie: #2.

Everyone worries about cheating and leaked data. Those are real risks, but they're governance problems. The bigger risk is deploying AI without a clear accountability contract. Before we deploy, I define what the system may do, what it may not do, and which decisions stay with a person. That's why our framework was built around human decision-making, not AI decision-making.

Five years in the title. This role asks for ten-plus, at principal level. Make the case.

Five years — I'll grant that. The posting's own "experience in lieu" clause is there for exactly this. But the title undercounts the work. Twenty years of high-stakes calls under strict rules came first, and then I rebuilt on purpose: degree, licence, engineer, product manager. Writing a governed AI framework — not just running one — is principal-level work, and it's already on its second production build. The bar counts years. The work shows judgment.

You say you "authored" the Operational AI Framework. Wasn't it handed to you?

I was handed the assignment: double call-review coverage, 5% to 10%. I wrote the framework: don't sample at all. Ingest, AI transform, human decision, feedback loop. Here's the proof. Enrollment coverage went from 5% to 100% at flat headcount. My business case put the hiring we avoided at about \$660K over five years. The customer's own estimate of the upside was about \$37M over five years.

Strip the story — isn't this just a QA efficiency tool?

No. It's one lever aimed down the whole funnel, end to end. Better feedback makes better front-line staff, and that moves what each stage owns: conversion, retention, and employment. One framework, three onboardings. Enrollment is the proof. Mentors is live. The rubric can be changed without retraining, and it has an exec champion and several business sponsors. Career Accelerators is in sea trial, on a Microsoft 365 engagement. Calibration went from 60% to 80%+, because every coaching review doubles as a training label.

AI scoring people's calls — why wasn't this resisted?

Because the trust contract was written at launch, not after someone pushed back. By definition, machine scores are non-determinative: they're never admissible in a performance review, and they only sequence which calls a human reviews. The models are in-house, inside the governance we already have, so there are no new data crossings and no procurement runway between a decision and a running system.

Name a product decision you owned that went wrong — and what it cost.

A leasability engine. It applied the Amazon discover-and-recommend pattern to leasing, back when nobody had Amazon's name on their lips. I built it unfunded, on my own time, with my best friend. Everyone I pitched it to dismissed it at first. Then they encouraged it. Nobody sponsored it. It made about \$2M in year one. I proved the loop was real by buying a hammock off an Instagram campaign inside a 1.5-mile walk, and we integrated a major retail partner in 30 days. It didn't die of product failure — it died because I mistook executive encouragement for executive sponsorship. That's the gate I now put on every build, and it's exactly the WGU design: in-house models, no new data crossings, no procurement runway. The scar is where the instinct comes from.

Law enforcement and a classroom. How does that translate to principal AI product?

Twenty years on the Priority of Life: hostages and victims first, then the public, then officers, then the subject. Risk is allocated in a fixed order, you decide without complete information, and you own the outcome. Three people are alive because I live inside that framework: a citizen, an officer, a subject. Ranking product work is the same discipline at lower stakes. It's why the framework keeps a human on the call that matters. Then a classroom, where students measured how high a rocket flew when they could never reach it — Alka-Seltzer launches, height by proxy, honest error bars. And I watched a girl get predicted to graduate on a passing grade she'd earned with a two-sentence paragraph across an entire term. A score that lies is worse than no score. That's the framework in miniature: score what you can't watch at scale, and say plainly what the score doesn't know.

Engineer to product manager in two years — how did that really happen?

The acquisition did it. AcceptanceNow became Acima, then Upbound. The finance work I'd built as an engineer — the automatic adjustments, the accounting method under Dodd-Frank inside and SOX outside, the Canadian tax logic — became the thing Accounting, AR/AP and both audit teams needed one person to answer for. I was already answering for it. The title caught up with the ownership. I didn't become a PM by asking. I became one by being the person the system already leaned on.

AND TWO QUESTIONS I WILL BE ASKING YOU

1. The first framework you'd want me to write here — not extend, write. What's the problem no one has taken up yet, and what has made it stick?
2. The framework got adopted because the trust contract — scores sequence, never evaluate — was written at launch. As Decision Intelligence scales agent-based systems across the lifecycle, where hasn't that contract been written yet? As a WGU student I keep landing on performance assessment rubrics, and the President and the Provost are asking for this, done right. Is that a direction you'd want brought inside?